

Improving Performance

Week 4

COMP6252 (Deep Learning Technologies)

ECS, University of Southampton

28 April 2022

Introduction

- As we have already seen DNN are powerful models.
- This is mainly due to their strong expressiveness:
 - the ability to model complicated relationships between input and output
- But this same expressiveness can lead to some problems

1 Overfitting

- Due to their power, NN can "learn" noise present in the training but not in the test datasets
- This reduces their generalization efficacy

2 Convergence

- Training takes long time. Sometimes doesn't even converge
- The results highly dependent on the meta parameters


What can be done?

This week we will introduce three approaches to mitigate the problems of convergence and overfitting.

- 1 Early stopping
- 2 Dropout
- 3 Batch normalization

Overfitting

- An example of overfitting is shown below
- The training accuracy keeps increasing while the validation plateau and even decreases




Early stopping

- In the previous figure, the max accuracy occurs at epoch 14
- If we stop training at epoch 14 or close to it not only we get better accuracy but we save lots of time
- This is exactly what is called **early stopping**
- How do we know when to stop training?
- One of the simplest methods is to monitor the **validation loss**


Validation loss vs accuracy

- Below is a plot of validation accuracy vs loss.
- Loss was multiplied by 0.5 to show them both on the same plot
- It is clear that accuracy plateaus after the loss starts increasing




A zoomed view

- Below is a zoomed view of the previous figure
- A vertical line passing by epoch 14 (where max accuracy occurs) is shown




Dropout

- Dropout is another method used to minimize overfitting
- A dropout layer "drops" certain nodes from the NN as shown in the figure below



(a) Standard Neural Net



(b) After applying dropout.

From: Srivastava et al.

What is happening exactly?

- Every training pass **different** nodes are dropped
- A neural network with n nodes that uses dropout, can be seen as a set of 2^n different "smaller" networks. Why?
- It means that using dropout is equivalent to averaging over an ensemble of Neural networks sharing the same weights
- Since the noise learned by each "smaller" network is random, "averaging" makes them cancel each other.

Implementation

- Using layers with an on/off switch on nodes is not practical
- More importantly, with an NN with large number of nodes it is **impossible**
- Instead, a node is turned off by multiplying its output with zero **with probability** p
- Mathematically, let \mathbf{y}^l be the output of layer l in a fully connected network
- Then the output of layer $(l + 1)$ is given by ($*$ is element wise product)

$$\mathbf{r}^l = \text{Bernoulli}(p)$$

$$\tilde{\mathbf{y}}^l = \mathbf{r}^l * \mathbf{y}^l$$

$$\mathbf{z}^{l+1} = \mathbf{W}^{l+1} \mathbf{y}^l + \mathbf{b}^{l+1}$$

$$\mathbf{y}^{l+1} = \sigma(\mathbf{z}^{l+1})$$

What about testing?

- We have seen that training a NN with dropout is equivalent to training multiple "smaller" networks.
- The result is an average over all such networks.
- To be consistent we must use the weights learned in the training phase for testing
- This means we have to **average** over all possible networks, which is infeasible.
- We approximate this averaging by **not using dropout** during testing
- But the weights obtained in the training phase are multiplied by the probability p
- Equivalently the weights are multiplied by $\frac{1}{p}$ during training

Batch Normalization

- A Batch normalization layer transforms the distribution of the input to a distribution that has 0 mean and unit variance
- Batch normalization is a powerful method that
 - 1 Allows for better convergence of NN training
 - 2 Adding BN layers to existing NN yield better generalization results
 - 3 Reduces over fitting
 - 4 Training is less sensitive to meta-parameters (learning rate, weights initialization)

Batch Normalization- How does it work

- Given a mini-batch of tensors x_{ci} of dimension (S,C,H,W)
- where c is the channel index and i collectively refers to all other dimensions.
- Let $N = S \times H \times W$.
- Batch normalization computes the mean and variance of the batch (per channel) according to

$$\mu_c = \frac{1}{N} \sum_{i=1}^N x_{ci}$$
$$\sigma_c^2 = \frac{1}{N} \sum_{i=1}^N (x_{ci} - \mu_c)^2$$

Batch Normalization- How does it work

The normalized inputs are computed as follows:

$$\hat{x}_{ci} = \frac{x_{ic} - \mu_c}{\sqrt{\sigma_c^2 + \epsilon}}$$

Therefore, for each channel, the \hat{x}_{ci} have zero mean and unit variance. The output of the batch normalization layer is given by

$$y_{ic} = \gamma \hat{x}_{ic} + \beta$$

Where γ and β are **learnable** parameters.

Batch normalization - Why does it work?

- Considerable ongoing debate
- In the original paper of [Ioffe and Szegedy](#) it claims that it fixes the internal covariate shift
- meaning the change in the distribution due to the change of the weights during learning
- [Kohler et al.](#) propose that it separates the change in magnitude from the change of directions.
- This can be seen from the fact that the \hat{x}_{ci} are normalized and all of them share the same magnitude γ

Example

Consider a mini-batch of size two, containing the tensors of size 2×2 , A and B with both having a single channel.

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

$$B = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$$

- The mean is 4.5 and (biased) variance is 5.25 therefore the output is

$$A = \begin{bmatrix} (1-4.5)/5.25 & (2-4.5)/5.25 \\ (3-4.5)/5.25 & (4-4.5)/5.25 \end{bmatrix}$$


$$B = \begin{bmatrix} (5-4.5)/5.25 & (6-4.5)/5.25 \\ (7-4.5)/5.25 & (8-4.5)/5.25 \end{bmatrix}$$

- When the input has multiple channels, the same operation is performed for each channel independently


In practice

- Data: CIFAR10
- The NN has two convolutional blocks
- The first contains two sub-blocks each with
 - 1 32 filters with receptive field of 3×3
 - 2 followed by BN layer
 - 3 followed by ReLU
- the sub-blocks are followed by a max pooling layer of size 2×2
- The second block is the same as the first but with 64 filter
- The conv blocks are followed by two feedforward layers with 128 and 10 nodes respectively

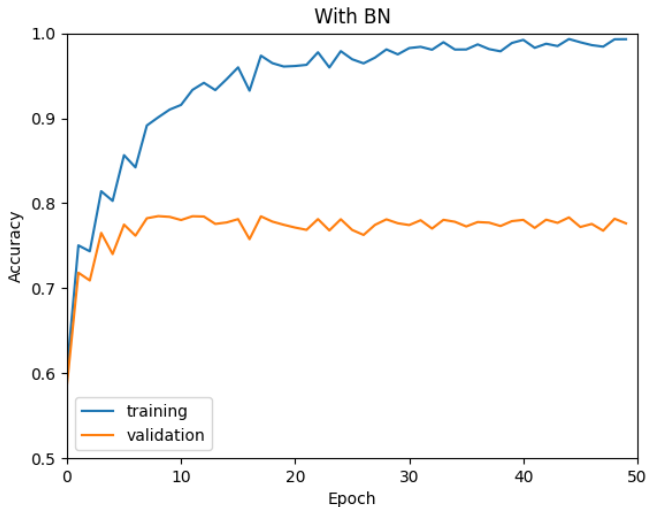
Network Architecture



Results: validation accuracy of BN vs no BN



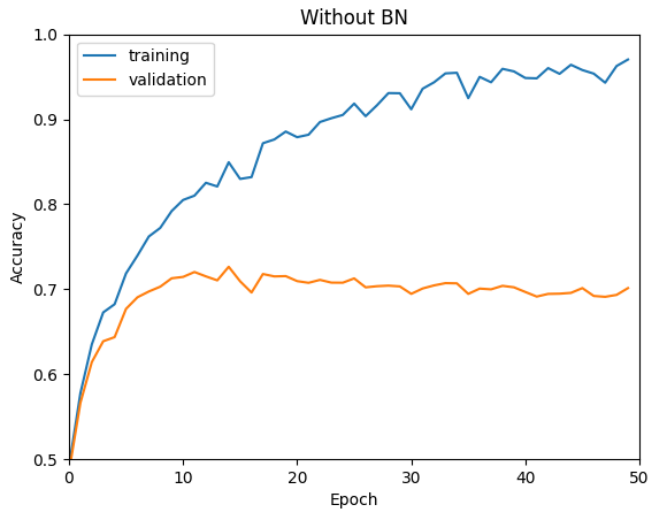
Results: train-vs-validation accuracy with BN






Analysis of results

- Batch normalization give much better accuracy.
- Peak value for the accuracy for BN is 78% vs 72%
- The average after the peak for BN is vs 75% vs 69%
- Also, with BN the peak acc is reached after 8 epochs vs 14 without BN.

Results: train-vs-validation accuracy without BN



References

-  Ioffe, Sergey and Christian Szegedy (2015). “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift”. In: *CoRR* abs/1502.03167. arXiv: 1502.03167. URL: <http://arxiv.org/abs/1502.03167>.
-  Kohler, Jonas et al. (2019). “Exponential convergence rates for batch normalization: The power of length-direction decoupling in non-convex optimization”. In: *The 22nd International Conference on Artificial Intelligence and Statistics*. PMLR, pp. 806–815.
-  Srivastava, Nitish et al. (2014). “Dropout: A Simple Way to Prevent Neural Networks from Overfitting”. In: *Journal of Machine Learning Research* 15.56, pp. 1929–1958. URL: <http://jmlr.org/papers/v15/srivastava14a.html>.